

# Step-by-Step Guide for Deployment on GCP with Correct File Paths

## 1. Upload Data and Scripts to GCS

Make sure you have already uploaded `movies.csv`, `ratings.csv`, and your PySpark script (e.g., `Recommendation_Engine_MovieLens.py`) to your GCS bucket as described earlier.

## 2. Create a Google Cloud Storage (GCS) Bucket

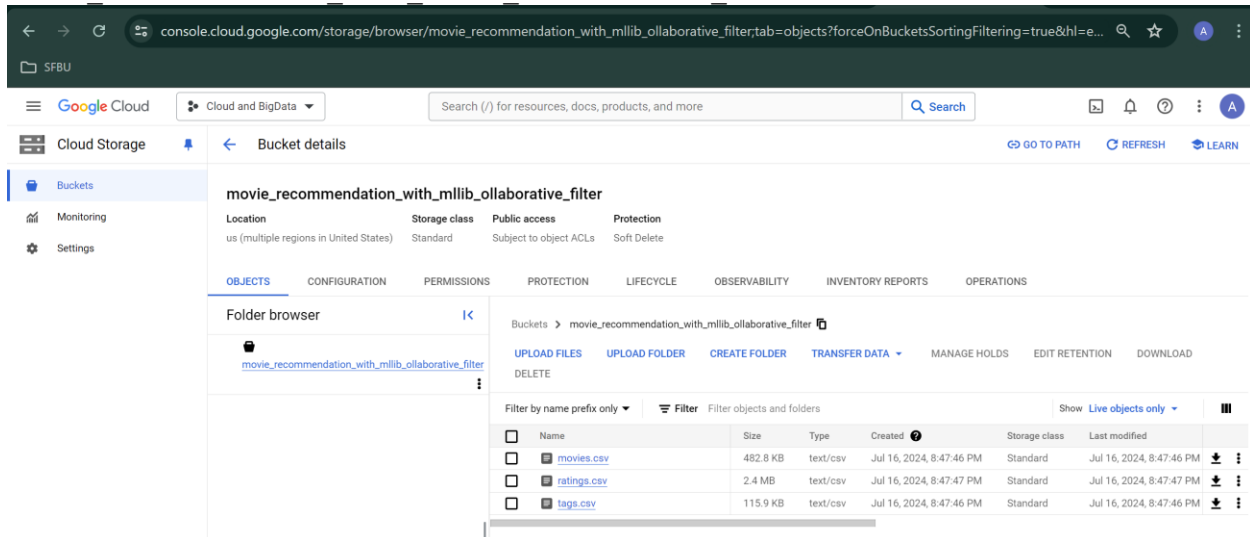
- Create a bucket in GCS to store your scripts and data.

```
gsutil mb gs://movie_recommendation_with_mllib_ollaborative_filter
adagniew407@cloudshell:~ (cloud-and-bigdata) $ gsutil mb gs://movie_recommendation_with_mllib_ollaborative_filter
Creating gs://movie_recommendation_with_mllib_ollaborative_filter/...
```

## 3. Upload Data and Scripts to GCS

- Upload the `movies.csv`, `ratings.csv`, and your PySpark script (e.g., `Recommendation_Engine_MovieLens.py`) to your GCS bucket.

```
gsutil cp movies.csv gs://
movie_recommendation_with_mllib_ollaborative_filter/
gsutil cp ratings.csv gs://
movie_recommendation_with_mllib_ollaborative_filter/
gsutil cp Recommendation_Engine_MovieLens.py gs://
movie_recommendation_with_mllib_ollaborative_filter
```



## 2. Modify the PySpark Script to Use GCS Paths

Update your PySpark script to read the files from GCS. You can use command-line arguments to pass the paths of the CSV files, making the script more flexible.

```
adagniew407@cloudshell:~ (cloud-and-bigdata) $ vim Recommendation_Engine_MovieLens.py
```

```
(train, test) = ratings.randomSplit([0.8, 0.2], seed=1234)

# Build ALS model
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating", nonnegative=True, implicitPrefs=False, coldStartStrategy="Drop")
param_grid = ParamGridBuilder() \
    .addGrid(als.rank, [10, 50, 100, 150]) \
    .addGrid(als.regParam, [.01, .05, .1, .15]) \
    .build()

evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating", predictionCol="prediction")
cv = CrossValidator(estimator=als, estimatorParamMaps=param_grid, evaluator=evaluator, numFolds=5)

# Train model
model = cv.fit(train)
best_model = model.bestModel

# Evaluate model
test_predictions = best_model.transform(test)
RMSE = evaluator.evaluate(test_predictions)
print(f"Root-mean-square error = (RMSE)")

# Generate recommendations
nrecommendations = best_model.recommendForAllUsers(10)
nrecommendations = nrecommendations \
    .withColumn("rec_exp", explode("recommendations")) \
    .select('userId', col("rec_exp.movieId"), col("rec_exp.rating"))
nrecommendations.show()

# Join with movie titles for better interpretability
nrecommendations.join(movies, on='movieId').filter('userId = 100').show()
ratings.join(movies, on='movieId').filter('userId = 100').sort('rating', ascending=False).limit(10).show()

# Stop Spark session
spark.stop()
```

Then, upload it to the bucket.

```
gsutil cp Recommendation_Engine_MovieLens.py
gs://movie_recommendation_with_mllib_ollaborative_filter
```

```
adagniew407@cloudshell:~ (cloud-and-bigdata)$ gsutil cp Recommendation_Engine_MovieLens.py gs://movie_recommendation_with_mllib_ollaborative_filter
Copying file://Recommendation_Engine_MovieLens.py [Content-Type=text/x-python]...
/ [1 files][ 2.2 KiB/ 2.2 KiB]
Operation completed over 1 objects/2.2 KiB.
adagniew407@cloudshell:~ (cloud-and-bigdata)$
```

Here is the script:

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import col, explode
from pyspark.ml.evaluation import RegressionEvaluator
from pyspark.ml.recommendation import ALS
from pyspark.ml.tuning import ParamGridBuilder, CrossValidator

import argparse
```

```

# Parse command-line arguments
parser = argparse.ArgumentParser()
parser.add_argument('--input_path_movies', required=True)
parser.add_argument('--input_path_ratings', required=True)
args = parser.parse_args()

# Initialize Spark session
spark = SparkSession.builder.appName('Recommendations').getOrCreate()

# Load data from GCS
movies = spark.read.csv(args.input_path_movies, header=True)
ratings = spark.read.csv(args.input_path_ratings, header=True)

# Preprocess data
ratings = ratings \
    .withColumn('userId', col('userId').cast('integer')) \
    .withColumn('movieId', col('movieId').cast('integer')) \
    .withColumn('rating', col('rating').cast('float')) \
    .drop('timestamp')

# Split data into training and testing sets
(train, test) = ratings.randomSplit([0.8, 0.2], seed=1234)

# Build ALS model
als = ALS(userCol="userId", itemCol="movieId", ratingCol="rating",
nonnegative=True, implicitPrefs=False, coldStartStrategy="drop")
param_grid = ParamGridBuilder() \
    .addGrid(als.rank, [10, 50, 100, 150]) \
    .addGrid(als.regParam, [.01, .05, .1, .15]) \
    .build()

evaluator = RegressionEvaluator(metricName="rmse", labelCol="rating",
predictionCol="prediction")
cv = CrossValidator(estimator=als, estimatorParamMaps=param_grid,
evaluator=evaluator, numFolds=5)

# Train model
model = cv.fit(train)
best_model = model.bestModel

# Evaluate model
test_predictions = best_model.transform(test)
RMSE = evaluator.evaluate(test_predictions)
print(f"Root-mean-square error = {RMSE}")

# Generate recommendations
nrecommendations = best_model.recommendForAllUsers(10)
nrecommendations = nrecommendations \
    .withColumn("rec_exp", explode("recommendations")) \
    .select('userId', col("rec_exp.movieId"), col("rec_exp.rating"))
nrecommendations.show()

# Join with movie titles for better interpretability
nrecommendations.join(movies, on='movieId').filter('userId = 100').show()
ratings.join(movies, on='movieId').filter('userId = 100').sort('rating',
ascending=False).limit(10).show()

```

```
# Stop Spark session
spark.stop()
```

### 3. create the cluster with the desired configuration:

```
gcloud dataproc clusters create spark-cluster \
  --region us-west1 \
  --zone us-west1-a \
  --master-machine-type n1-standard-4 \
  --worker-machine-type n1-standard-4 \
  --num-workers 2

adagniew407@cloudshell:~ (cloud-and-bigdata)$ gcloud dataproc clusters create spark-cluster \
  --region us-west1 \
  --zone us-west1-a \
  --master-machine-type n1-standard-4 \
  --worker-machine-type n1-standard-4 \
  --num-workers 2

Waiting on operation [projects/cloud-and-bigdata/regions/us-west1/operations/0435e684-2e21-360f-b3b8-70bcd09be86].
Waiting for cluster creation operation...
WARNING: No image specified. Using the default image version. It is recommended to select a specific image version in production, as the default image version may change at any time.
WARNING: Consider using Auto Zone rather than selecting a zone manually. See https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/auto-zone
WARNING: Failed to validate permissions required for default service account: '635972213485-compute@developer.gserviceaccount.com'. Cluster creation could still be successful if required permissions have been granted to the respective service accounts as mentioned in the document https://cloud.google.com/dataproc/docs/concepts/configuring-clusters/service-accounts#dataproc\_service\_accounts\_2. This could be due to Cloud Resource Manager API hasn't been enabled in your project '635972213485' before or it is disabled. Enable it by visiting 'https://console.developers.google.com/apis/api/cloudresourcemanager.googleapis.com/overview?project=635972213485'.
WARNING: The firewall rules for specified network or subnetwork would allow ingress traffic from 0.0.0.0/0, which could be a security risk.
WARNING: The specified custom staging bucket 'dataproc-staging-us-west1-635972213485-dblsca5v' is not using uniform bucket level access IAM configuration. It is recommended to update bucket to enable the same. See https://cloud.google.com/storage/docs/uniform-bucket-level-access.
Waiting for cluster creation operation...working...
```

### 4. Submit the PySpark Job with GCS Paths

Submit your PySpark job to the Dataproc cluster, providing the GCS paths for the input files:

```
gcloud dataproc jobs submit pyspark gs://your-bucket-name/Recommendation_Engine_MovieLens.py \
  --cluster=spark-cluster \
  --region=us-west1 \
  -- \
  --input_path_movies=gs://movie_recommendation_with_mllib_ollaborative_filter/movies.csv \
  --input_path_ratings=gs://movie_recommendation_with_mllib_ollaborative_filter/ratings.csv
```

Replace your-bucket-name with the actual name of your GCS bucket.

```
adagniew407@cloudshell:~ (cloud-and-bigdata)$ gcloud dataproc jobs submit pyspark gs://movie_recommendation_with_mllib_ollaborative_filter/Recommendation_Engine_MovieLens.py --cluster=spark-cluster --region=us-west1 --input_path_movies=gs://movie_recommendation_with_mllib_ollaborative_filter/movies.csv --input_path_ratings=gs://movie_recommendation_with_mllib_ollaborative_filter/ratings.csv
Job [710caa605793471089b145cf98731e15] submitted.
Waiting for job output...
24/07/17 04:00:51 INFO org.apache.spark.SparkEnv: Registering MapOutputTracker
24/07/17 04:00:51 INFO org.apache.spark.SparkEnv: Registering BlockManagerMaster
24/07/17 04:00:51 INFO org.apache.spark.SparkEnv: Registering BlockManagerMasterHeartbeat
24/07/17 04:00:51 INFO org.apache.spark.SparkEnv: Registering OutputCommitCoordinator
24/07/17 04:00:51 INFO org.sparkproject.jetty.util.log: Logging initialized @5607ms to org.sparkproject.jetty.util.log.Slf4jLog
24/07/17 04:00:51 INFO org.sparkproject.jetty.server.Server: jetty-9.4.40.v20210413; built: 2021-04-13T20:42:42.668Z; git: b881a572662e1943a14ae12e7e1207989f210b74; jvm 1.8.0_412-b08
```

Root-mean-square error = 0.8685666272031686

userId	movieId	rating
540	3379	5.4218884
540	33649	5.060796
540	171495	5.0543323
540	5490	4.947019
540	179135	4.9249244
540	26073	4.9249244
540	7071	4.9249244
540	84273	4.9249244
540	184245	4.9249244
540	117531	4.9249244
580	3379	4.814103
580	33649	4.7172403
580	6300	4.7001023
580	171495	4.6310377
580	179135	4.6122727
580	117531	4.6122727
580	84273	4.6122727
580	7071	4.6122727
580	184245	4.6122727
580	26073	4.6122727

only showing top 20 rows

```

+-----+-----+-----+-----+-----+
|movieId|userId|  rating|          title|          genres|
+-----+-----+-----+-----+-----+
|  67618|   100|5.1201425|Strictly Sexual (...|Comedy|Drama|Romance|
|   3379|   100| 5.064743|On the Beach (1959)|          Drama|
|  42730|   100| 5.042285|  Glory Road (2006)|          Drama|
|  33649|   100| 5.021657|  Saving Face (2004)|Comedy|Drama|Romance|
| 117531|   100|4.9267745|  Watermark (2014)|          Documentary|
|   7071|   100|4.9267745|Woman Under the I...|          Drama|
| 184245|   100|4.9267745|De platte jungle ...|          Documentary|
|   26073|   100|4.9267745|Human Condition I...|          Drama|War|
| 179135|   100|4.9267745|Blue Planet II (2...|          Documentary|
|   84273|   100|4.9267745|Zeitgeist: Moving...|          Documentary|
+-----+-----+-----+-----+-----+

+-----+-----+-----+-----+-----+
|movieId|userId|rating|          title|          genres|
+-----+-----+-----+-----+-----+
|   1101|   100|   5.0|   Top Gun (1986)|   Action|Romance| |
|   1958|   100|   5.0|Terms of Endearme...|   Comedy|Drama|
|   2423|   100|   5.0|Christmas Vacatio...|   Comedy|
|   4041|   100|   5.0|Officer and a Gen...|   Drama|Romance|
|   5620|   100|   5.0|Sweet Home Alabam...|   Comedy|Romance|
|    368|   100|   4.5|   Maverick (1994)|Adventure|Comedy|...|
|    934|   100|   4.5|Father of the Bri...|   Comedy|
|    539|   100|   4.5|Sleepless in Seat...|Comedy|Drama|Romance|
|     16|   100|   4.5|   Casino (1995)|   Crime|Drama|
|    553|   100|   4.5|   Tombstone (1993)|Action|Drama|Western|
+-----+-----+-----+-----+-----+

```

By following these steps, your PySpark script will correctly read the files from GCS when running on GCP Dataproc.